

Class 7: Machine Learning

Courtney Anderson (PID:A69038035)

Today, we will begin our exploration of some "classical" means learning approaches. We will start with clustering:

Lets first make up some data cluster where we know what the answer should be.

```
hist (rnorm(1000))
```

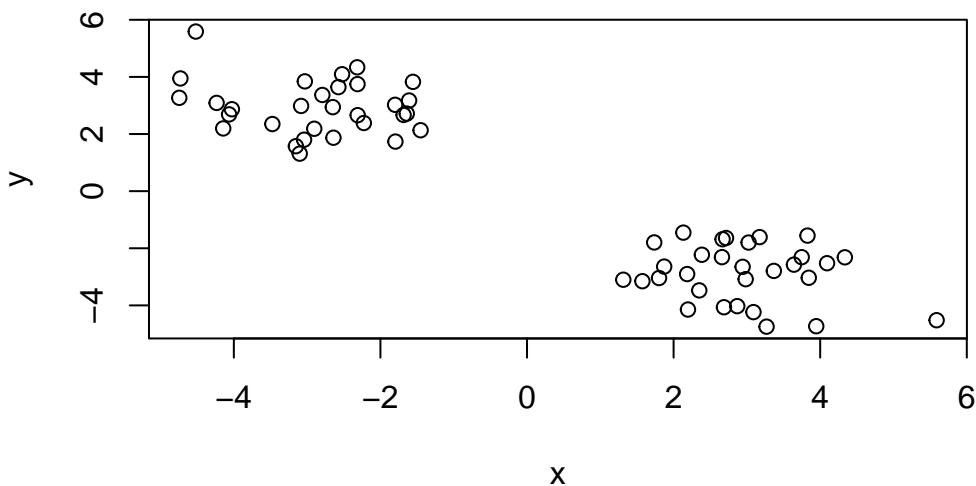


```
x <- c( rnorm(30, mean=3), rnorm(30,mean=-3))
y <- rev(x)

z <- cbind(x,y)
head(z)
```

```
x           y
[1,] 1.801183 -3.041708
[2,] 2.941477 -2.649415
[3,] 2.184707 -2.902584
[4,] 3.826722 -1.556109
[5,] 3.946130 -4.729025
[6,] 2.715266 -1.640327
```

```
plot(z)
```



The main function in “base” R for K-means clustering is called `kmeans()`.

```
k <- kmeans(z, centers = 2)
```

Q. How big are the clusters?

```
k$size
```

```
[1] 30 30
```

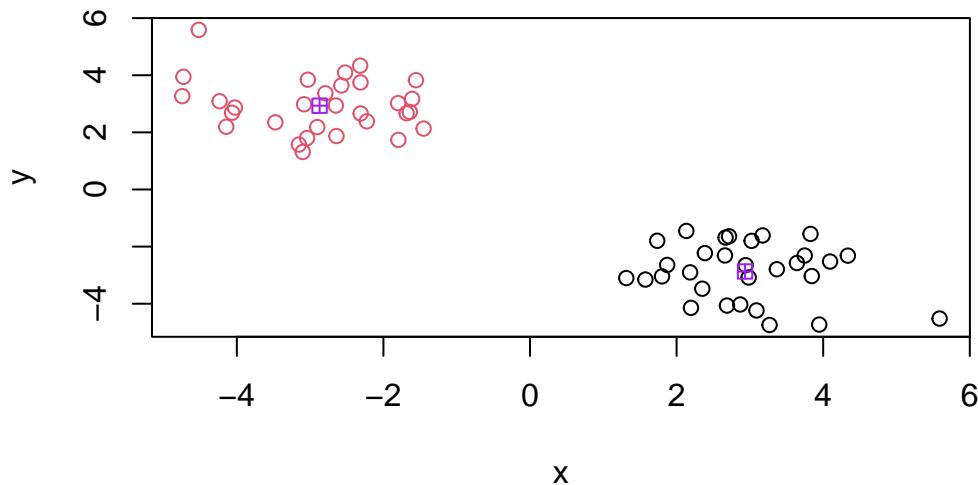
Q. What clusters do my data points reside in?

```
k$cluster
```

```
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1  
[39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

Q. can you make a plot of our data colored by cluster assignment i.e. Make a result figure

```
plot(z, col = k$cluster)  
points(k$centers, col = "purple", pch = 12)
```



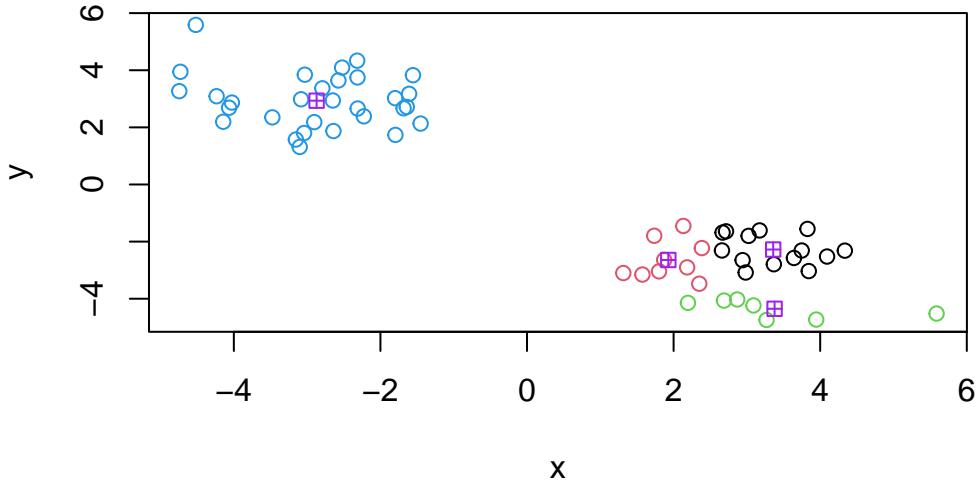
Q. Cluster with K-means into 4 clusters and plot your results above

```
k4 <- kmeans(z,centers = 4)
```

```
k4$cluster
```

```
[1] 2 1 2 1 3 1 2 3 3 1 3 1 1 1 1 1 2 2 2 1 1 1 2 3 1 3 3 2 1 4 4 4 4 4 4  
[39] 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4
```

```
plot(z, col=k4$cluster)
points(k4$centers, col="purple", pch=12)
```



Run kmeans with centers (i.e. values of k) equal 1 to 6

```
k1 <- kmeans(x, centers=1)$tot.witness
k2 <- kmeans(x, centers=2)$tot.witness
k3 <- kmeans(x, centers=3)$tot.witness
k4 <- kmeans(x, centers=4)$tot.witness
k5 <- kmeans(x, centers=5)$tot.witness
k6 <- kmeans(x, centers=6)$tot.witness

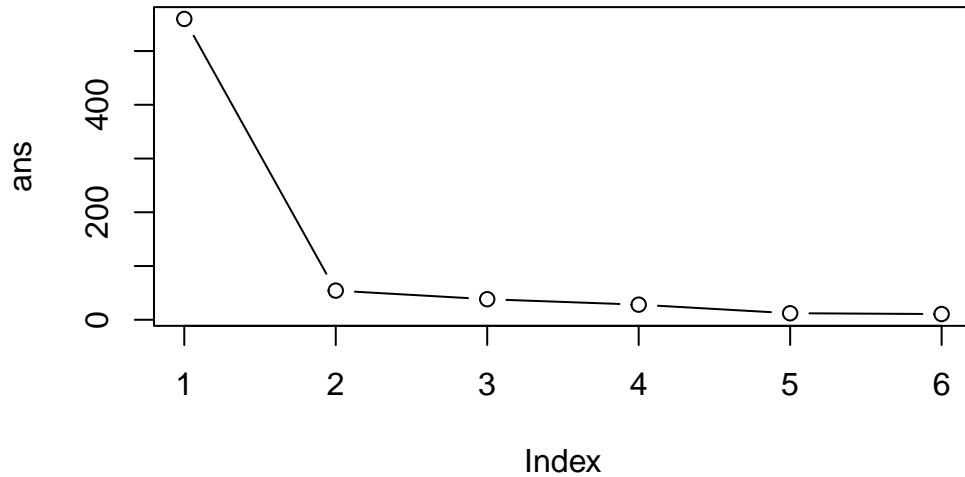
ans <- c(k1, k2, k3, k4, k5, k6)
```

Or use a for loop (Make a scree-plot)

```
ans <- NULL
for(i in 1:6){
  ans <- c(ans, kmeans(x, centers=i)$tot.withinss)
}
ans
```

```
[1] 559.73038 54.28315 38.27174 28.14872 12.13730 10.66589
```

```
plot(ans, typ="b")
```



```
##Hierarchical Clustering
```

The main function in “base” R for this is called `hclust()`

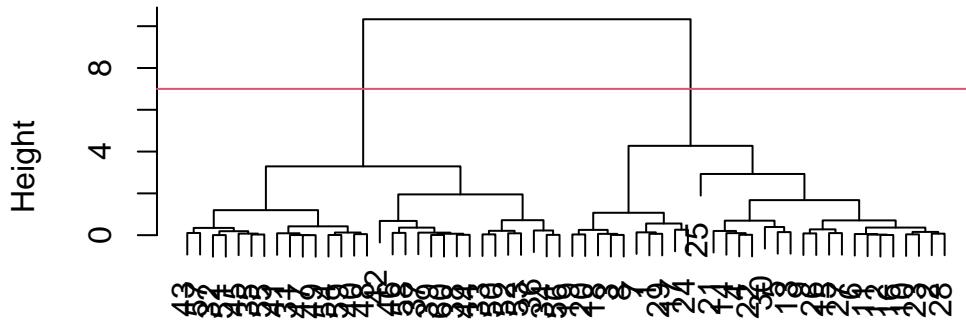
```
d <- dist(x)
hc <- hclust(d)
hc
```

```
Call:
hclust(d = d)
```

```
Cluster method : complete
Distance       : euclidean
Number of objects: 60
```

```
plot(hc)
abline(h=7,col=2)
```

Cluster Dendrogram



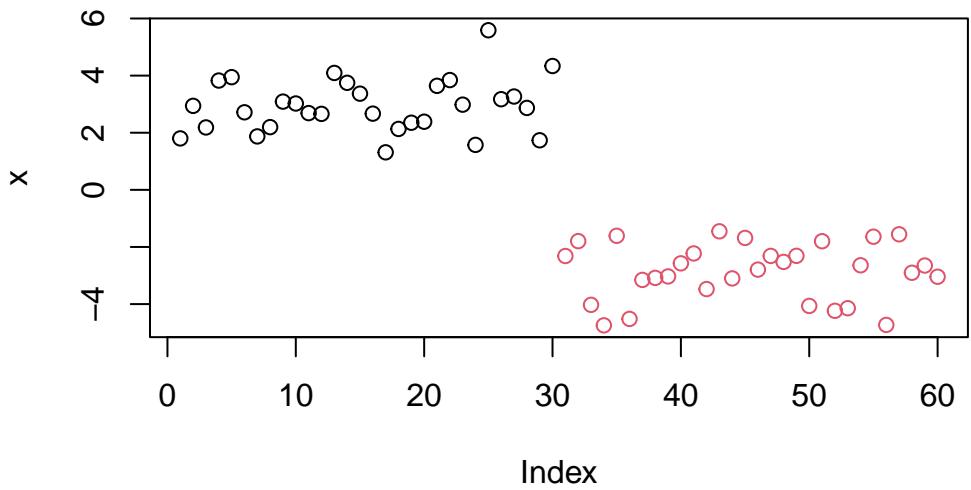
```
d  
hclust (*, "complete")
```

To obtain clusters from our `hclust` object `hc` we “cut” the tree into sub branches. For this we can use the `cutree()` function

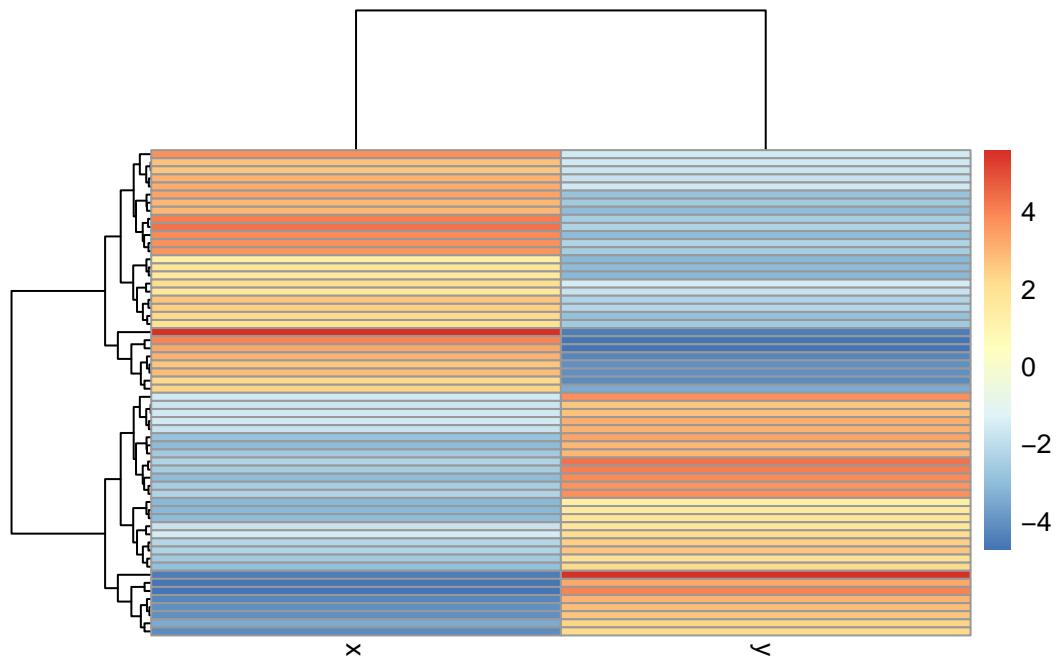
```
grps <- cutree(hc, h=7)  
grps
```

Results figure

```
plot(x,col=grps)
```



```
library (pheatmap)
pheatmap(z)
```



Principal Component Analysis (PCA)

PCA is a dimensional reduction, to take all things measuring and projects them on PC axes. PC1 is the “best fit” of the data, that maximizes the data spread/variance in data. It captures the most variance. PC2 captures the rest of the variance. PC looks to see the left/right and up/down variance easier to see. Think of a big funnel, and putting all your datta in it.

Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

There are 17 rows and 5 columns. You could use nrow and ncol.

```
x <- read.csv(file = "UK_foods.csv")
```

```
nrow(x)
```

```
[1] 17
```

```
ncol(x)
```

```
[1] 5
```

Preview the first 6 rows

```
head(x)
```

	X	England	Wales	Scotland	N.Ireland
1	Cheese	105	103	103	66
2	Carcass_meat	245	227	242	267
3	Other_meat	685	803	750	586
4	Fish	147	160	122	93
5	Fats_and_oils	193	235	184	209
6	Sugars	156	175	147	139

Fix the row names

```
rownames(x) <- x[,1]
x <- x[,-1]
head(x)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

```
dim(x)
```

[1] 17 4

I fixed the rownames again (I was having trouble with my file)

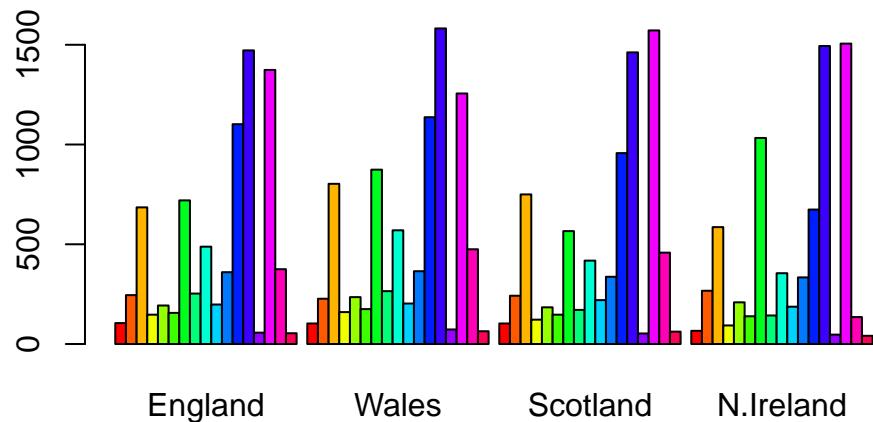
```
x <- read.csv("UK_foods.csv", row.names=1)
head(x)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

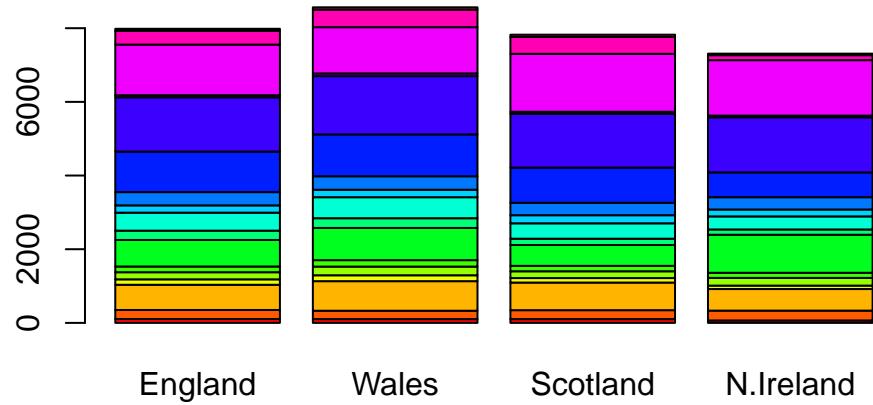
Q2. Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

I prefer the second method just because its quicker and requires less typing. (more)

```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```



```
barplot(as.matrix(x), beside=F, col=rainbow(nrow(x)))
```



Q3: Changing what optional argument in the above barplot() function results in the following plot?

Just change the beside to equal False or F

```
library(tidyr)
x_long <- x |>

tibble::rownames_to_column("Food") |>
pivot_longer(cols = -Food,
names_to = "Country",
values_to = "Consumption")

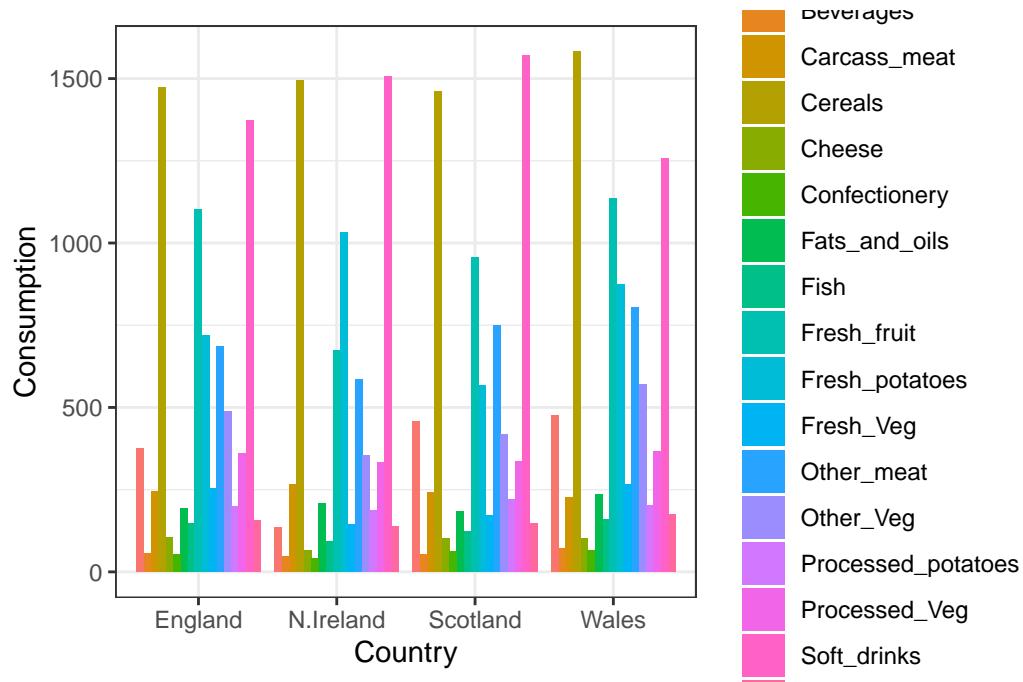
dim(x_long)
```

```
[1] 68 3
```

```
head(x_long)
```

```
# A tibble: 6 x 3
  Food           Country   Consumption
  <chr>          <chr>        <int>
1 "Cheese"       England      105
2 "Cheese"       Wales        103
3 "Cheese"       Scotland     103
4 "Cheese"       N.Ireland    66
5 "Carcass_meat" England     245
6 "Carcass_meat" Wales       227
```

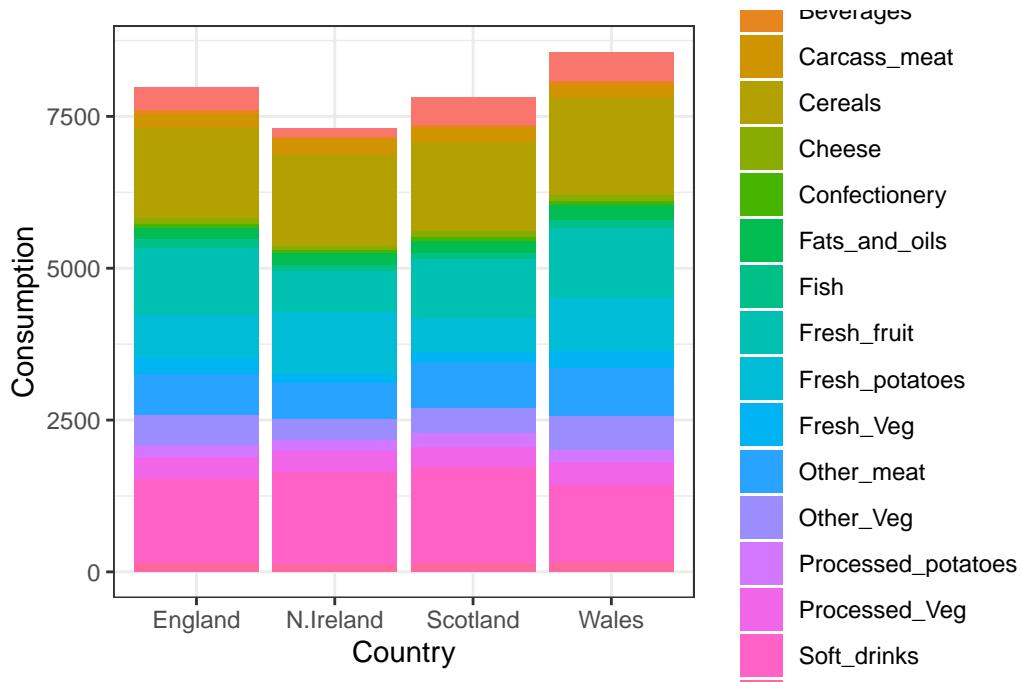
```
library(ggplot2)
ggplot(x_long) +
aes(Country, Consumption, fill = Food) +
geom_col(position = "dodge") +
theme_bw()
```



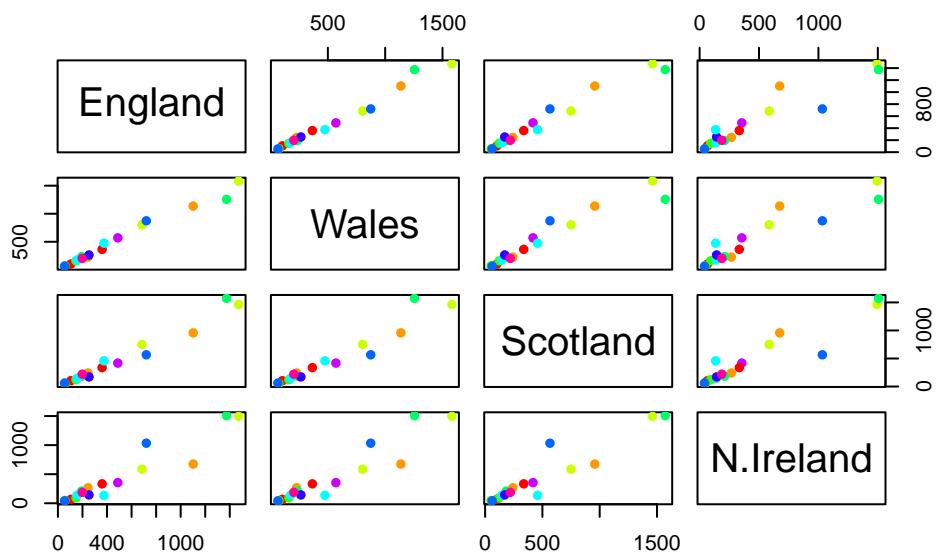
Q4: Changing what optional argument in the above barplot() function results in the following plot?

Change the geom_col(position = "dodge") to just geom_col().

```
ggplot(x_long) +
  aes(Country, Consumption, fill = Food) +
  geom_col() +
  theme_bw()
```



```
pairs(x, col=rainbow(10), pch=16)
```

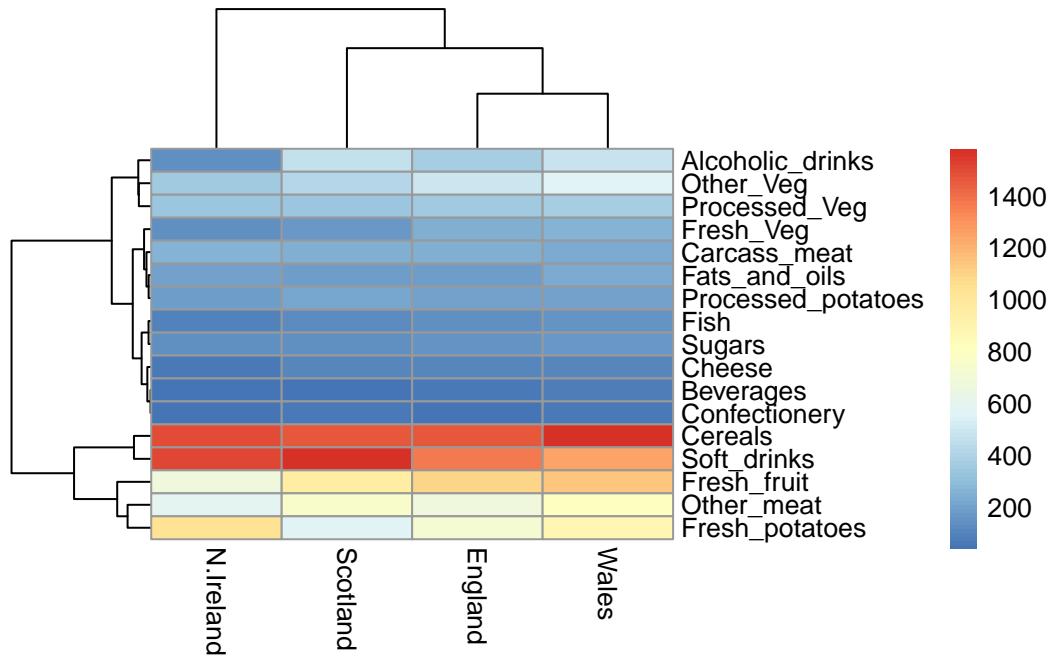


Q5: Generating all pairwise plots may help somewhat. Can you make sense of the

following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

The points on the plots are specific types of food. If a point lies on the diagonal then that means that the food level is the same in both countries. If one point lies outside of the diagonal then that means that the food is being consumed at different levels between the two countries.

```
pheatmap( as.matrix(x) )
```



Q6. What is the main differences between N. Ireland and the other countries of the UK in terms of this data-set?

Ireland appears to consume more potatoes than the other countries. They also consume less fresh fruit and other meats compared to other countries. Finally they have high consumption of soft drinks and cereals (but these values are more comparable).

```
##PCA to the rescue
```

Q7. Complete the code below to generate a plot of PC1 vs PC2. The second line adds text labels over the data points.

The main function in “base” R for PCA is called `prcomp()`. As we want to do PCA on the food data, for the different countries, we will want the foods in the columns.

```
t(x)
```

	Cheese	Carcass_meat	Other_meat	Fish	Fats_and_oils	Sugars
England	105	245	685	147	193	156
Wales	103	227	803	160	235	175
Scotland	103	242	750	122	184	147
N.Ireland	66	267	586	93	209	139
	Fresh_potatoes	Fresh_Veg	Other_Veg	Processed_potatoes		
England	720	253	488		198	
Wales	874	265	570		203	
Scotland	566	171	418		220	
N.Ireland	1033	143	355		187	
	Processed_Veg	Fresh_fruit	Cereals	Beverages	Soft_drinks	
England	360	1102	1472	57	1374	
Wales	365	1137	1582	73	1256	
Scotland	337	957	1462	53	1572	
N.Ireland	334	674	1494	47	1506	
	Alcoholic_drinks	Confectionery				
England	375	54				
Wales	475	64				
Scotland	458	62				
N.Ireland	135	41				

```
pca <- prcomp( t(x) )
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	2.921e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

```
pca$x
```

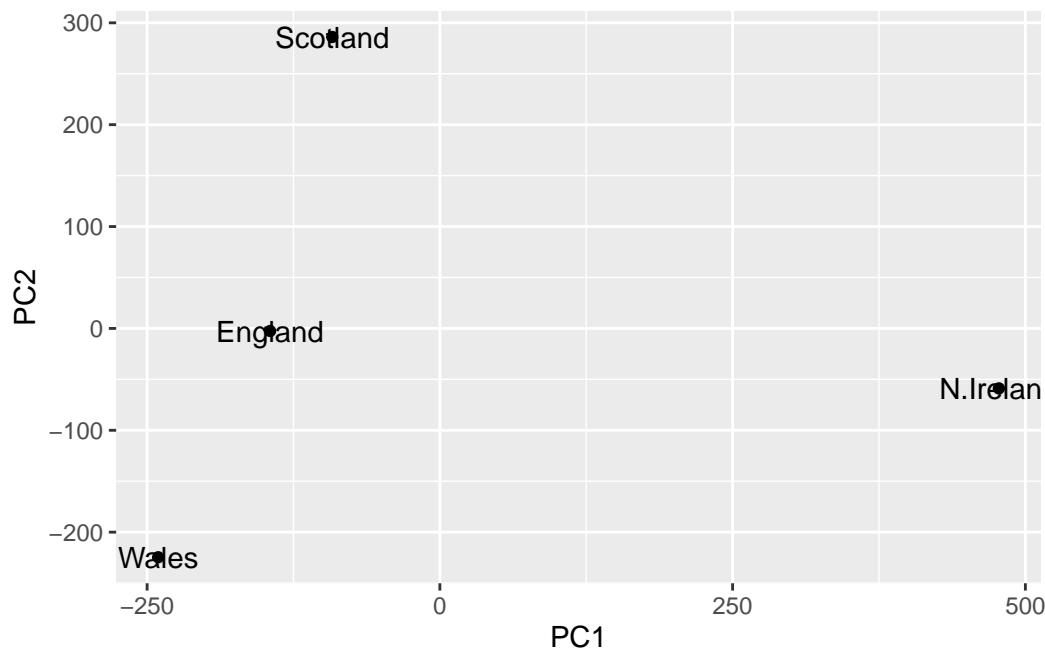
	PC1	PC2	PC3	PC4
England	-144.99315	-2.532999	105.768945	-9.152022e-15
Wales	-240.52915	-224.646925	-56.475555	5.560040e-13
Scotland	-91.86934	286.081786	-44.415495	-6.638419e-13
N.Ireland	477.39164	-58.901862	-4.877895	1.329771e-13

```

library(ggplot2)

ggplot(pca$x) +
  aes(PC1, PC2, label = rownames(pca$x)) +
  geom_point() +
  geom_text()

```



Q8. Customize your plot so that the colors of the country names match the colors in our UK and Ireland map and table at start of this document.

```

library(ggplot2)

# Example: add a grouping column (one color per region)
regions <- data.frame(
  name = rownames(pca$x),
  region = c("England", "Scotland", "Wales", "N.Ireland")
)

# Combine with PCA results
df <- cbind(pca$x, regions)

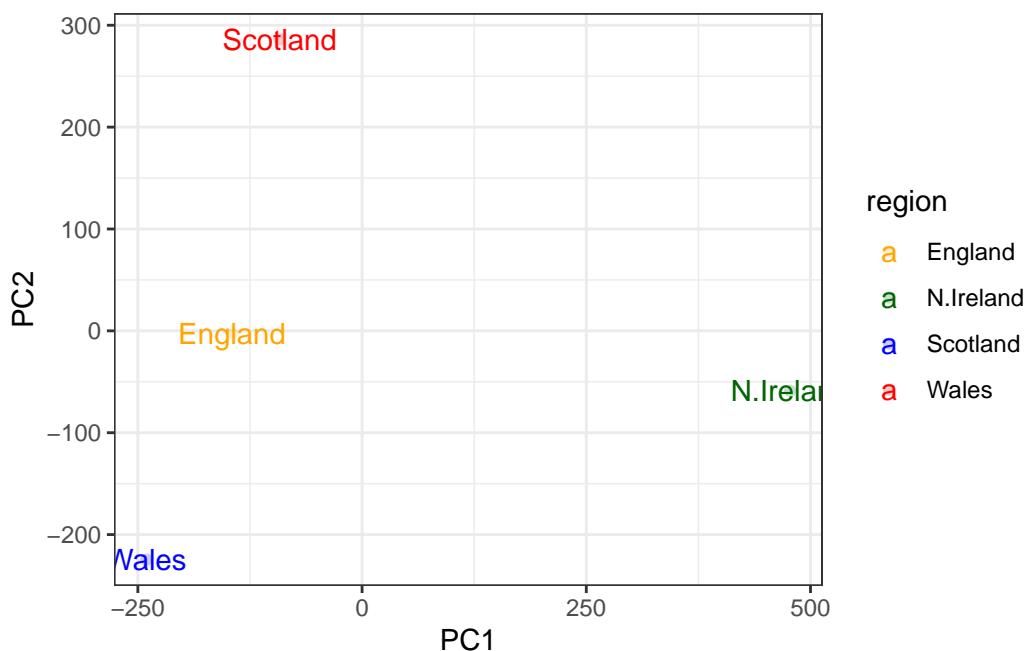
# Plot with custom colors

```

```

ggplot(df, aes(PC1, PC2, label = name, color = region)) +
  geom_point(alpha=0.2) +
  geom_text() +
  scale_color_manual(values = c(
    "England" = "orange",
    "Scotland" = "blue",
    "Wales" = "red",
    "N.Ireland" = "darkgreen"
  )) +
  theme_bw() +
  labs(x = "PC1", y = "PC2")

```



```

v <- round( pca$sdev^2/sum(pca$sdev^2) * 100 )
v

```

```
[1] 67 29 4 0
```

```

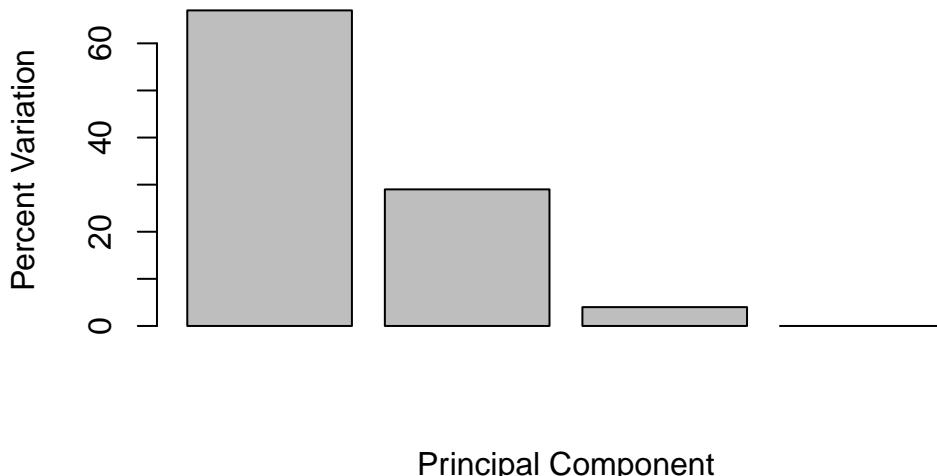
z <- summary(pca)
z$importance

```

PC1	PC2	PC3	PC4
-----	-----	-----	-----

```
Standard deviation      324.15019 212.74780 73.87622 2.921348e-14
Proportion of Variance  0.67444   0.29052   0.03503  0.000000e+00
Cumulative Proportion   0.67444   0.96497   1.00000  1.000000e+00
```

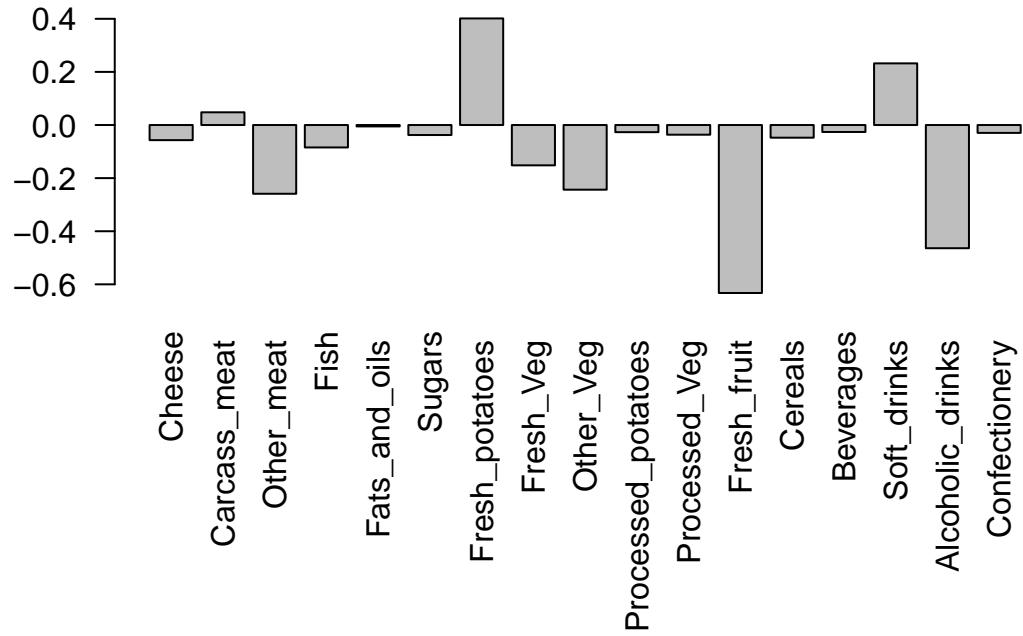
```
barplot(v, xlab="Principal Component", ylab="Percent Variation")
```



```
##Digging Deeper(variable loadings)
```

Lets focus on PC1 as it accounts for > 90% of variance

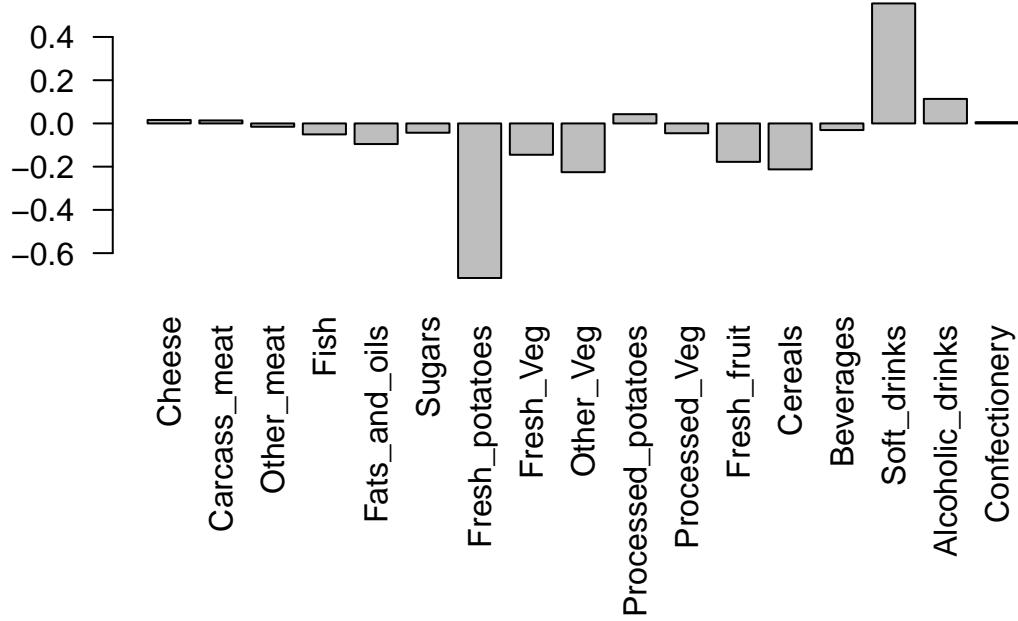
```
par(mar=c(10, 3, 0.35, 0))
barplot( pca$rotation[,1], las=2 )
```



Q9: Generate a similar ‘loadings plot’ for PC2. What two food groups feature prominently and what does PC2 mainly tell us about?

Fresh potatoes and soft drinks are featured prominently. PC2 mainly tells us the secondary differences in the data.

```
par(mar=c(10, 3, 0.35, 0))
barplot( pca$rotation[,2], las=2 )
```



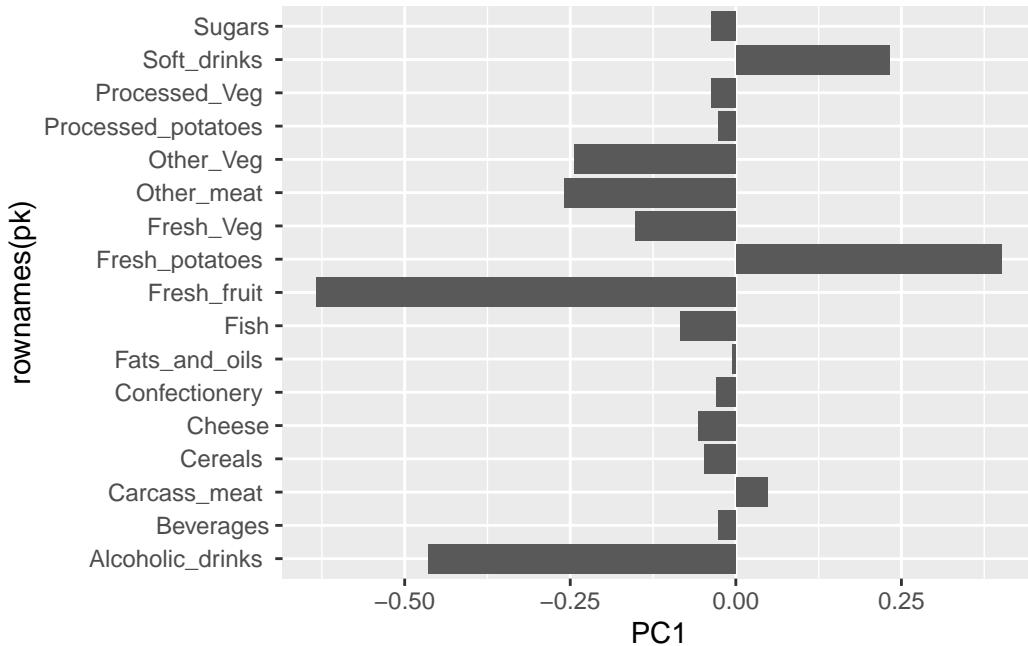
Another major result of PCA is the so-called “variable loadings” or \$rotation that tells us how the original variables (foods) contribute to PCs (i.e. our new axis).

```
pk <- pca$rotation
pk
```

	PC1	PC2	PC3	PC4
Cheese	-0.056955380	0.016012850	0.02394295	-0.409382587
Carcass_meat	0.047927628	0.013915823	0.06367111	0.729481922
Other_meat	-0.258916658	-0.015331138	-0.55384854	0.331001134
Fish	-0.084414983	-0.050754947	0.03906481	0.022375878
Fats_and_oils	-0.005193623	-0.095388656	-0.12522257	0.034512161
Sugars	-0.037620983	-0.043021699	-0.03605745	0.024943337
Fresh_potatoes	0.401402060	-0.715017078	-0.20668248	0.021396007
Fresh_Veg	-0.151849942	-0.144900268	0.21382237	0.001606882
Other_Veg	-0.243593729	-0.225450923	-0.05332841	0.031153231
Processed_potatoes	-0.026886233	0.042850761	-0.07364902	-0.017379680
Processed_Veg	-0.036488269	-0.045451802	0.05289191	0.021250980
Fresh_fruit	-0.632640898	-0.177740743	0.40012865	0.227657348
Cereals	-0.047702858	-0.212599678	-0.35884921	0.100043319
Beverages	-0.026187756	-0.030560542	-0.04135860	-0.018382072
Soft_drinks	0.232244140	0.555124311	-0.16942648	0.222319484
Alcoholic_drinks	-0.463968168	0.113536523	-0.49858320	-0.273126013

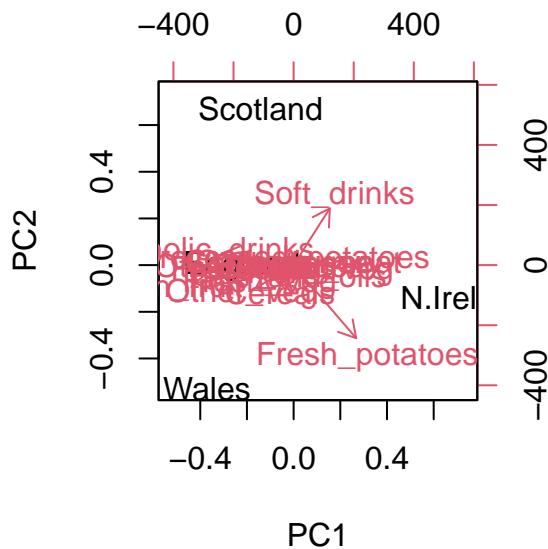
```
Confectionery      -0.029650201  0.005949921 -0.05232164  0.001890737
```

```
ggplot(pk) + aes(PC1, rownames(pk)) + geom_col()
```



```
##Biplots
```

```
biplot(pca)
```



##PCA or RNA seq data

```
url2 <- "https://tinyurl.com/expression-CSV"
rna.data <- read.csv(url2, row.names=1)
head(rna.data)
```

	wt1	wt2	wt3	wt4	wt5	ko1	ko2	ko3	ko4	ko5
gene1	439	458	408	429	420	90	88	86	90	93
gene2	219	200	204	210	187	427	423	434	433	426
gene3	1006	989	1030	1017	973	252	237	238	226	210
gene4	783	792	829	856	760	849	856	835	885	894
gene5	181	249	204	244	225	277	305	272	270	279
gene6	460	502	491	491	493	612	594	577	618	638

```
nrow(rna.data)
```

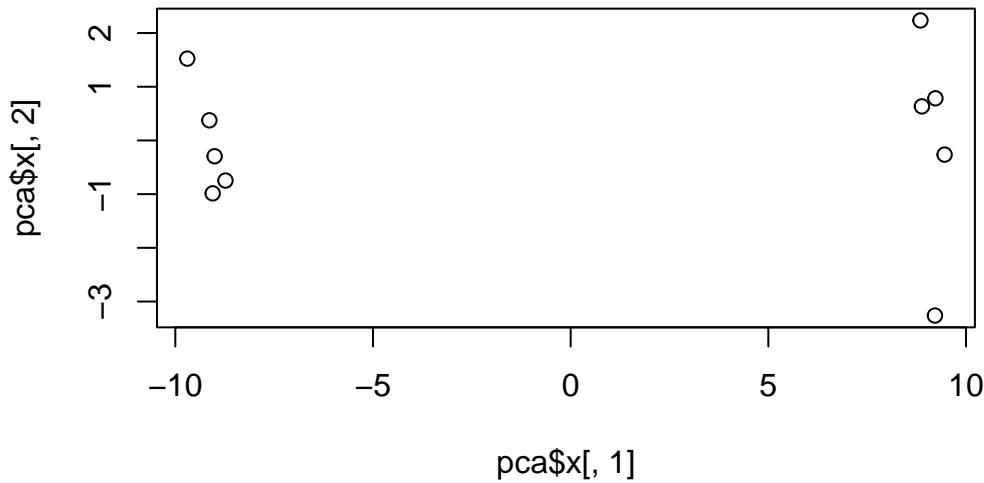
[1] 100

Q10: How many genes and samples are in this data set?

There are 100 genes.

```
pca <- prcomp(t(rna.data), scale=TRUE)
```

```
plot(pca$x[,1], pca$x[,2])
```



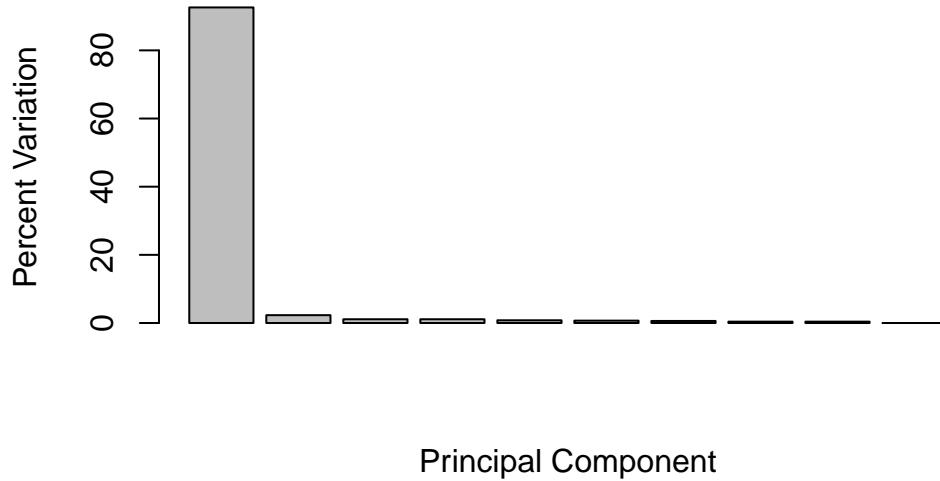
```
pca.var <- pca$sdev^2
```

```
pca.var.per <- round(pca.var/sum(pca.var)*100, 1)
pca.var.per
```

```
[1] 92.6 2.3 1.1 1.1 0.8 0.7 0.6 0.4 0.4 0.0
```

```
barplot(pca.var.per, main="Scree Plot",
        xlab="Principal Component", ylab="Percent Variation")
```

Scree Plot

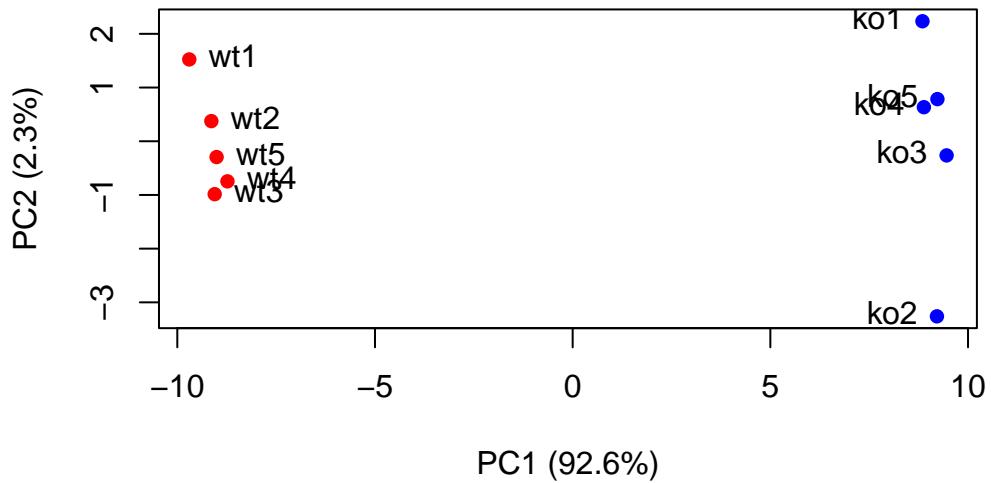


A vector of colors for wt and ko samples

```
colvec <- colnames(rna.data)
colvec[grep("wt", colvec)] <- "red"
colvec[grep("ko", colvec)] <- "blue"

plot(pca$x[,1], pca$x[,2], col=colvec, pch=16,
      xlab=paste0("PC1 (", pca.var.per[1], "%)"),
      ylab=paste0("PC2 (", pca.var.per[2], "%)"))

text(pca$x[,1], pca$x[,2], labels = colnames(rna.data), pos=c(rep(4,5), rep(2,5)))
```



Another way to color by sample type

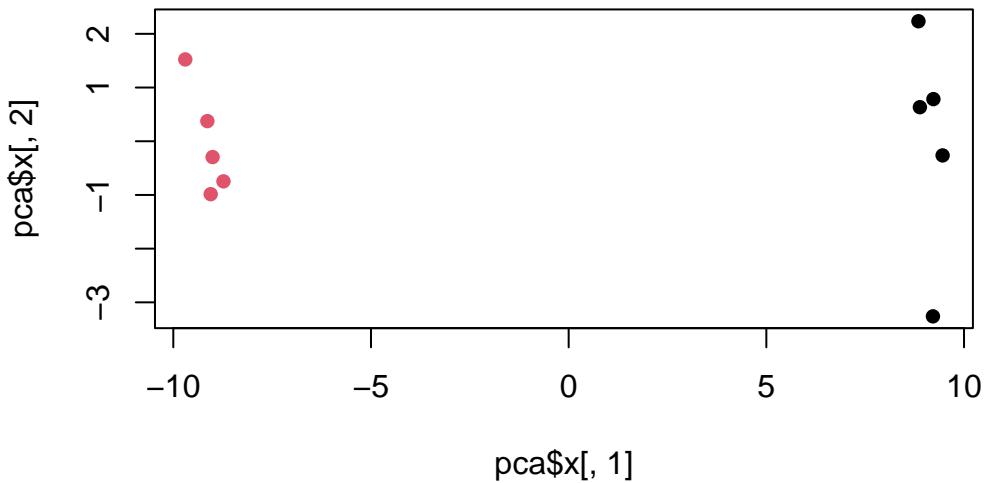
Extract the first 2 characters of the sample name

```
sample.type <- substr(colnames(rna.data), 1, 2)
sample.type
```

```
[1] "wt" "wt" "wt" "wt" "wt" "ko" "ko" "ko" "ko" "ko"
```

now use this as a factor input to color our plot

```
plot(pca$x[, 1], pca$x[, 2], col=as.factor(sample.type), pch=16)
```



**Find the top 10 measurements (genes) that contribute
most to PC1 in either direction (+ or -)**

```
loading_scores <- pca$rotation[,1]
gene_scores <- abs(loading_scores)
gene_score_ranked <- sort(gene_scores, decreasing=TRUE)
```

Show the top ten genes

```
top_10_genes <- names(gene_score_ranked[1:10])
top_10_genes
```

```
[1] "gene100" "gene66"  "gene45"   "gene68"   "gene98"   "gene60"   "gene21"
[8] "gene56"   "gene10"   "gene90"
```