

Class 8: Mini Project

Courtney Anderson (PID:A69038035)

Table of contents

| | |
|--|----|
| Background | 1 |
| Data Import | 2 |
| Exploring my data | 3 |
| Principal Component Analysis | 6 |
| Performing PCA | 6 |
| Interpreting PCA Results | 8 |
| Variance Explained: Making Scree Plots | 11 |
| Communicating PCA results | 13 |
| Hierarchical Clustering | 14 |
| Results of Hierarchical Clustering | 14 |
| Selecting number of clusters | 15 |
| Using Different Methods | 15 |
| Combining methods | 16 |
| Create a Ward clustering | 16 |
| Sensitivity/Specificity | 20 |
| Prediction | 20 |

Background

The goal of this mini-project is for you to explore a complete analysis using the unsupervised learning techniques covered in class. You'll extend what you've learned by combining PCA as a preprocessing step to clustering using data that consist of measurements of cell nuclei of human breast masses. This expands on our RNA-Seq analysis from last day.

The data itself comes from the Wisconsin Breast Cancer Diagnostic Data Set first reported by K. P. Benne and O. L. Mangasarian: "Robust Linear Programming Discrimination of Two Linearly Inseparable Sets".

Values in this data set describe characteristics of the cell nuclei present in digitized images of a fine needle aspiration (FNA) of a breast mass.

Data Import

Read in data and name it; edit the roles:

```
fna.data <- "WisconsinCancer.csv"  
wisc.df <- read.csv(fna.data, row.names=1)  
head(wisc.df)
```

| | diagnosis | radius_mean | texture_mean | perimeter_mean | area_mean |
|----------|-----------------|------------------------|----------------|---------------------|-------------------|
| 842302 | M | 17.99 | 10.38 | 122.80 | 1001.0 |
| 842517 | M | 20.57 | 17.77 | 132.90 | 1326.0 |
| 84300903 | M | 19.69 | 21.25 | 130.00 | 1203.0 |
| 84348301 | M | 11.42 | 20.38 | 77.58 | 386.1 |
| 84358402 | M | 20.29 | 14.34 | 135.10 | 1297.0 |
| 843786 | M | 12.45 | 15.70 | 82.57 | 477.1 |
| | smoothness_mean | compactness_mean | concavity_mean | concave.points_mean | |
| 842302 | 0.11840 | 0.27760 | 0.3001 | | 0.14710 |
| 842517 | 0.08474 | 0.07864 | 0.0869 | | 0.07017 |
| 84300903 | 0.10960 | 0.15990 | 0.1974 | | 0.12790 |
| 84348301 | 0.14250 | 0.28390 | 0.2414 | | 0.10520 |
| 84358402 | 0.10030 | 0.13280 | 0.1980 | | 0.10430 |
| 843786 | 0.12780 | 0.17000 | 0.1578 | | 0.08089 |
| | symmetry_mean | fractal_dimension_mean | radius_se | texture_se | perimeter_se |
| 842302 | 0.2419 | | 0.07871 | 1.0950 | 0.9053 |
| 842517 | 0.1812 | | 0.05667 | 0.5435 | 0.7339 |
| 84300903 | 0.2069 | | 0.05999 | 0.7456 | 0.7869 |
| 84348301 | 0.2597 | | 0.09744 | 0.4956 | 1.1560 |
| 84358402 | 0.1809 | | 0.05883 | 0.7572 | 0.7813 |
| 843786 | 0.2087 | | 0.07613 | 0.3345 | 0.8902 |
| | area_se | smoothness_se | compactness_se | concavity_se | concave.points_se |
| 842302 | 153.40 | 0.006399 | 0.04904 | 0.05373 | 0.01587 |
| 842517 | 74.08 | 0.005225 | 0.01308 | 0.01860 | 0.01340 |
| 84300903 | 94.03 | 0.006150 | 0.04006 | 0.03832 | 0.02058 |
| 84348301 | 27.23 | 0.009110 | 0.07458 | 0.05661 | 0.01867 |
| 84358402 | 94.44 | 0.011490 | 0.02461 | 0.05688 | 0.01885 |
| 843786 | 27.19 | 0.007510 | 0.03345 | 0.03672 | 0.01137 |
| | symmetry_se | fractal_dimension_se | radius_worst | texture_worst | |
| 842302 | 0.03003 | 0.006193 | 25.38 | 17.33 | |
| 842517 | 0.01389 | 0.003532 | 24.99 | 23.41 | |
| 84300903 | 0.02250 | 0.004571 | 23.57 | 25.53 | |
| 84348301 | 0.05963 | 0.009208 | 14.91 | 26.50 | |
| 84358402 | 0.01756 | 0.005115 | 22.54 | 16.67 | |

| | | | | |
|----------|-------------------------|----------------------|------------------|-------------------|
| 843786 | 0.02165 | 0.005082 | 15.47 | 23.75 |
| | perimeter_worst | area_worst | smoothness_worst | compactness_worst |
| 842302 | 184.60 | 2019.0 | 0.1622 | 0.6656 |
| 842517 | 158.80 | 1956.0 | 0.1238 | 0.1866 |
| 84300903 | 152.50 | 1709.0 | 0.1444 | 0.4245 |
| 84348301 | 98.87 | 567.7 | 0.2098 | 0.8663 |
| 84358402 | 152.20 | 1575.0 | 0.1374 | 0.2050 |
| 843786 | 103.40 | 741.6 | 0.1791 | 0.5249 |
| | concavity_worst | concave.points_worst | symmetry_worst | |
| 842302 | 0.7119 | 0.2654 | 0.4601 | |
| 842517 | 0.2416 | 0.1860 | 0.2750 | |
| 84300903 | 0.4504 | 0.2430 | 0.3613 | |
| 84348301 | 0.6869 | 0.2575 | 0.6638 | |
| 84358402 | 0.4000 | 0.1625 | 0.2364 | |
| 843786 | 0.5355 | 0.1741 | 0.3985 | |
| | fractal_dimension_worst | | | |
| 842302 | 0.11890 | | | |
| 842517 | 0.08902 | | | |
| 84300903 | 0.08758 | | | |
| 84348301 | 0.17300 | | | |
| 84358402 | 0.07678 | | | |
| 843786 | 0.12440 | | | |

Remove the first column; we dont want to add patient IDs:

```
wisc.data <- wisc.df[,-1]
```

Create a diagnosis vector:

```
diagnosis <- as.factor(wisc.df$diagnosis)
```

Exploring my data

```
head(wisc.data)
```

| | radius_mean | texture_mean | perimeter_mean | area_mean | smoothness_mean |
|----------|-------------|--------------|----------------|-----------|-----------------|
| 842302 | 17.99 | 10.38 | 122.80 | 1001.0 | 0.11840 |
| 842517 | 20.57 | 17.77 | 132.90 | 1326.0 | 0.08474 |
| 84300903 | 19.69 | 21.25 | 130.00 | 1203.0 | 0.10960 |

| | | | | | |
|---|----------|----------|---------|---------|---------|
| 84348301 | 11.42 | 20.38 | 77.58 | 386.1 | 0.14250 |
| 84358402 | 20.29 | 14.34 | 135.10 | 1297.0 | 0.10030 |
| 843786 | 12.45 | 15.70 | 82.57 | 477.1 | 0.12780 |
| compactness_mean concavity_mean concave.points_mean symmetry_mean | | | | | |
| 842302 | 0.27760 | 0.3001 | 0.14710 | 0.2419 | |
| 842517 | 0.07864 | 0.0869 | 0.07017 | 0.1812 | |
| 84300903 | 0.15990 | 0.1974 | 0.12790 | 0.2069 | |
| 84348301 | 0.28390 | 0.2414 | 0.10520 | 0.2597 | |
| 84358402 | 0.13280 | 0.1980 | 0.10430 | 0.1809 | |
| 843786 | 0.17000 | 0.1578 | 0.08089 | 0.2087 | |
| fractal_dimension_mean radius_se texture_se perimeter_se area_se | | | | | |
| 842302 | 0.07871 | 1.0950 | 0.9053 | 8.589 | 153.40 |
| 842517 | 0.05667 | 0.5435 | 0.7339 | 3.398 | 74.08 |
| 84300903 | 0.05999 | 0.7456 | 0.7869 | 4.585 | 94.03 |
| 84348301 | 0.09744 | 0.4956 | 1.1560 | 3.445 | 27.23 |
| 84358402 | 0.05883 | 0.7572 | 0.7813 | 5.438 | 94.44 |
| 843786 | 0.07613 | 0.3345 | 0.8902 | 2.217 | 27.19 |
| smoothness_se compactness_se concavity_se concave.points_se | | | | | |
| 842302 | 0.006399 | 0.04904 | 0.05373 | 0.01587 | |
| 842517 | 0.005225 | 0.01308 | 0.01860 | 0.01340 | |
| 84300903 | 0.006150 | 0.04006 | 0.03832 | 0.02058 | |
| 84348301 | 0.009110 | 0.07458 | 0.05661 | 0.01867 | |
| 84358402 | 0.011490 | 0.02461 | 0.05688 | 0.01885 | |
| 843786 | 0.007510 | 0.03345 | 0.03672 | 0.01137 | |
| symmetry_se fractal_dimension_se radius_worst texture_worst | | | | | |
| 842302 | 0.03003 | 0.006193 | 25.38 | 17.33 | |
| 842517 | 0.01389 | 0.003532 | 24.99 | 23.41 | |
| 84300903 | 0.02250 | 0.004571 | 23.57 | 25.53 | |
| 84348301 | 0.05963 | 0.009208 | 14.91 | 26.50 | |
| 84358402 | 0.01756 | 0.005115 | 22.54 | 16.67 | |
| 843786 | 0.02165 | 0.005082 | 15.47 | 23.75 | |
| perimeter_worst area_worst smoothness_worst compactness_worst | | | | | |
| 842302 | 184.60 | 2019.0 | 0.1622 | 0.6656 | |
| 842517 | 158.80 | 1956.0 | 0.1238 | 0.1866 | |
| 84300903 | 152.50 | 1709.0 | 0.1444 | 0.4245 | |
| 84348301 | 98.87 | 567.7 | 0.2098 | 0.8663 | |
| 84358402 | 152.20 | 1575.0 | 0.1374 | 0.2050 | |
| 843786 | 103.40 | 741.6 | 0.1791 | 0.5249 | |
| concavity_worst concave.points_worst symmetry_worst | | | | | |
| 842302 | 0.7119 | 0.2654 | 0.4601 | | |
| 842517 | 0.2416 | 0.1860 | 0.2750 | | |
| 84300903 | 0.4504 | 0.2430 | 0.3613 | | |
| 84348301 | 0.6869 | 0.2575 | 0.6638 | | |

```
84358402          0.4000      0.1625      0.2364
843786          0.5355      0.1741      0.3985
fractal_dimension_worst
842302          0.11890
842517          0.08902
84300903      0.08758
84348301      0.17300
84358402      0.07678
843786          0.12440
```

```
head(diagnosis)
```

```
[1] M M M M M M  
Levels: B M
```

Q1. How many observations are in this dataset?

```
nrow(wisc.data)
```

```
[1] 569
```

There are 569 observations in this data set.

Q2. How many of the observations have a malignant diagnosis?

```
table(diagnosis)
```

```
diagnosis
B     M
357 212
```

212 observations have been a malignant diagnosis.

Q3. How many variables/features in the data are suffixed with _mean?

```
length(grep("_mean$", names(wisc.data)))
```

```
[1] 10
```

10 variables are suffixed with _mean.

Principal Component Analysis

The main function in base R for PCA is called `prcomp()`. An optional argument `scal` should nearly always be switched to `scale=TRUE` for this function.

Performing PCA

Check column means and standard deviations:

```
colMeans(wisc.data)
```

| | | |
|------------------------|----------------------|-------------------------|
| radius_mean | texture_mean | perimeter_mean |
| 1.412729e+01 | 1.928965e+01 | 9.196903e+01 |
| area_mean | smoothness_mean | compactness_mean |
| 6.548891e+02 | 9.636028e-02 | 1.043410e-01 |
| concavity_mean | concave.points_mean | symmetry_mean |
| 8.879932e-02 | 4.891915e-02 | 1.811619e-01 |
| fractal_dimension_mean | radius_se | texture_se |
| 6.279761e-02 | 4.051721e-01 | 1.216853e+00 |
| perimeter_se | area_se | smoothness_se |
| 2.866059e+00 | 4.033708e+01 | 7.040979e-03 |
| compactness_se | concavity_se | concave.points_se |
| 2.547814e-02 | 3.189372e-02 | 1.179614e-02 |
| symmetry_se | fractal_dimension_se | radius_worst |
| 2.054230e-02 | 3.794904e-03 | 1.626919e+01 |
| texture_worst | perimeter_worst | area_worst |
| 2.567722e+01 | 1.072612e+02 | 8.805831e+02 |
| smoothness_worst | compactness_worst | concavity_worst |
| 1.323686e-01 | 2.542650e-01 | 2.721885e-01 |
| concave.points_worst | symmetry_worst | fractal_dimension_worst |
| 1.146062e-01 | 2.900756e-01 | 8.394582e-02 |

```
apply(wisc.data, 2, sd)
```

| | | |
|----------------|---------------------|------------------|
| radius_mean | texture_mean | perimeter_mean |
| 3.524049e+00 | 4.301036e+00 | 2.429898e+01 |
| area_mean | smoothness_mean | compactness_mean |
| 3.519141e+02 | 1.406413e-02 | 5.281276e-02 |
| concavity_mean | concave.points_mean | symmetry_mean |
| 7.971981e-02 | 3.880284e-02 | 2.741428e-02 |

| | | | |
|------------------------|----------------------|-------------------------|-------------------|
| fractal_dimension_mean | | radius_se | texture_se |
| 7.060363e-03 | | 2.773127e-01 | 5.516484e-01 |
| perimeter_se | | area_se | smoothness_se |
| 2.021855e+00 | | 4.549101e+01 | 3.002518e-03 |
| compactness_se | | concavity_se | concave.points_se |
| 1.790818e-02 | | 3.018606e-02 | 6.170285e-03 |
| symmetry_se | fractal_dimension_se | | radius_worst |
| 8.266372e-03 | | 2.646071e-03 | 4.833242e+00 |
| texture_worst | perimeter_worst | | area_worst |
| 6.146258e+00 | | 3.360254e+01 | 5.693570e+02 |
| smoothness_worst | compactness_worst | concavity_worst | |
| 2.283243e-02 | | 1.573365e-01 | 2.086243e-01 |
| concave.points_worst | symmetry_worst | fractal_dimension_worst | |
| 6.573234e-02 | | 6.186747e-02 | 1.806127e-02 |

Perform PCA and inspect summary:

```
wisc.pr <- prcomp(wisc.data, scale=T )
summary (wisc.pr)
```

Importance of components:

| | PC1 | PC2 | PC3 | PC4 | PC5 | PC6 | PC7 |
|------------------------|---------|---------|---------|---------|---------|---------|---------|
| Standard deviation | 3.6444 | 2.3857 | 1.67867 | 1.40735 | 1.28403 | 1.09880 | 0.82172 |
| Proportion of Variance | 0.4427 | 0.1897 | 0.09393 | 0.06602 | 0.05496 | 0.04025 | 0.02251 |
| Cumulative Proportion | 0.4427 | 0.6324 | 0.72636 | 0.79239 | 0.84734 | 0.88759 | 0.91010 |
| | PC8 | PC9 | PC10 | PC11 | PC12 | PC13 | PC14 |
| Standard deviation | 0.69037 | 0.6457 | 0.59219 | 0.5421 | 0.51104 | 0.49128 | 0.39624 |
| Proportion of Variance | 0.01589 | 0.0139 | 0.01169 | 0.0098 | 0.00871 | 0.00805 | 0.00523 |
| Cumulative Proportion | 0.92598 | 0.9399 | 0.95157 | 0.9614 | 0.97007 | 0.97812 | 0.98335 |
| | PC15 | PC16 | PC17 | PC18 | PC19 | PC20 | PC21 |
| Standard deviation | 0.30681 | 0.28260 | 0.24372 | 0.22939 | 0.22244 | 0.17652 | 0.1731 |
| Proportion of Variance | 0.00314 | 0.00266 | 0.00198 | 0.00175 | 0.00165 | 0.00104 | 0.0010 |
| Cumulative Proportion | 0.98649 | 0.98915 | 0.99113 | 0.99288 | 0.99453 | 0.99557 | 0.9966 |
| | PC22 | PC23 | PC24 | PC25 | PC26 | PC27 | PC28 |
| Standard deviation | 0.16565 | 0.15602 | 0.1344 | 0.12442 | 0.09043 | 0.08307 | 0.03987 |
| Proportion of Variance | 0.00091 | 0.00081 | 0.0006 | 0.00052 | 0.00027 | 0.00023 | 0.00005 |
| Cumulative Proportion | 0.99749 | 0.99830 | 0.9989 | 0.99942 | 0.99969 | 0.99992 | 0.99997 |
| | PC29 | PC30 | | | | | |
| Standard deviation | 0.02736 | 0.01153 | | | | | |
| Proportion of Variance | 0.00002 | 0.00000 | | | | | |
| Cumulative Proportion | 1.00000 | 1.00000 | | | | | |

Q4. From your results, what proportion of the original variance is captured by the first principal components (PC1)? 44% of the original variance is captured by PC1.


Q5. How many principal components (PCs) are required to describe at least 70% of the original variance in the data? Three principle components would be required to describe at least 70% of the original variance of the data.

Q6. How many principal components (PCs) are required to describe at least 90% of the original variance in the data? Seven principle components would be required to describe at least 90% of the original variance of the data.

Interpreting PCA Results

Create a biplot:

```
biplot(wisc.pr)
```



Q7. What stands out to you about this plot? Is it easy or difficult to understand? Why? This plot is very difficult to understand because everything is overlapping.


Create a scatterplot:

```

library(ggplot2)

ggplot(wisc.pr$x, aes(x = PC1, y = PC2, color = diagnosis)) +
  geom_point(size = 2) +
  labs(x = "PC1", y = "PC2", title = "PC1 vs PC2 colored by diagnosis") +
  theme_minimal()

```




Create a similar plot for components 1 and 3:

```

ggplot(wisc.pr$x, aes(x = PC1, y = PC3, color = diagnosis)) +
  geom_point(size = 2) +
  labs(x = "PC1", y = "PC2", title = "PC1 vs PC2 colored by diagnosis") +
  theme_minimal()

```




Q8. Generate a similar plot for principal components 1 and 3. What do you notice about these plots? In the second plot, the points appear to be shifted downward.

Create plots using ggplot:

```
df <- as.data.frame(wisc.pr$x)
df$diagnosis <- diagnosis
```

```
library(ggplot2)
ggplot(df) +
  aes(PC1, PC2, col=diagnosis) +
  geom_point()
```



Variance Explained: Making Scree Plots


Calculate the variance of each component:

```
pr.var <- wisc.pr$sdev^2
head(pr.var)
```

```
[1] 13.281608 5.691355 2.817949 1.980640 1.648731 1.207357
```


Explain and plot variance:

```
pve <- wisc.pr$sdev^2 / sum(wisc.pr$sdev^2)
plot(pve, xlab = "Principal Component",
     ylab = "Proportion of Variance Explained",
     ylim = c(0, 1), type = "o")
```



Create an alternate scree plot:

```
barplot(pve, ylab = "Percent of Variance Explained",
        names.arg=paste0("PC",1:length(pve)), las=2, axes = FALSE)
axis(2, at=pve, labels=round(pve,2)*100 )
```



Communicating PCA results

Q9. For the first principal component, what is the component of the loading vector (i.e. `wisc.pr$rotation[,1]`) for the feature `concave.points_mean`? This tells us how much this original feature contributes to the first PC.

```
wisc.pr$rotation[,1]
```

| | | |
|------------------------|----------------------|-------------------|
| radius_mean | texture_mean | perimeter_mean |
| -0.21890244 | -0.10372458 | -0.22753729 |
| area_mean | smoothness_mean | compactness_mean |
| -0.22099499 | -0.14258969 | -0.23928535 |
| concavity_mean | concave.points_mean | symmetry_mean |
| -0.25840048 | -0.26085376 | -0.13816696 |
| fractal_dimension_mean | radius_se | texture_se |
| -0.06436335 | -0.20597878 | -0.01742803 |
| perimeter_se | area_se | smoothness_se |
| -0.21132592 | -0.20286964 | -0.01453145 |
| compactness_se | concavity_se | concave.points_se |
| -0.17039345 | -0.15358979 | -0.18341740 |
| symmetry_se | fractal_dimension_se | radius_worst |

| | | |
|----------------------|-------------------|-------------------------|
| -0.04249842 | -0.10256832 | -0.22799663 |
| texture_worst | perimeter_worst | area_worst |
| -0.10446933 | -0.23663968 | -0.22487053 |
| smoothness_worst | compactness_worst | concavity_worst |
| -0.12795256 | -0.21009588 | -0.22876753 |
| concave.points_worst | symmetry_worst | fractal_dimension_worst |
| -0.25088597 | -0.12290456 | -0.13178394 |

The feature for concave points mean is -0.26085376.

Hierarchical Clustering

Scale the original data:

```
data.scaled <- scale(wisc.data)
```

Calculate the distances between all pairs of observations in the new scaled dataset:

```
data.dist <- dist(data.scaled, method= "euclidean")
```

Create a hierarchical clustering using complete linkage:


```
wisc.hclust <- hclust(data.dist, method="complete")
```

Results of Hierarchical Clustering

Q10. Using the plot() and abline() functions, what is the height at which the clustering model has 4 clusters?

```
plot(wisc.hclust)
abline(h = wisc.hclust$height[length(wisc.hclust$height) - 3], col = "red")
```

Cluster Dendrogram



```
data.dist  
hclust (*, "complete")
```

Selecting number of clusters

Use cutree to cut the tree so that it has 4 clusters; show a table:

```
wisc.hclust.clusters <- cutree(wisc.hclust, k = 4)  
table(wisc.hclust.clusters, diagnosis)
```

| wisc.hclust.clusters | diagnosis | |
|----------------------|-----------|-----|
| | B | M |
| 1 | 12 | 165 |
| 2 | 2 | 5 |
| 3 | 343 | 40 |
| 4 | 0 | 2 |

Q11. OPTIONAL: Can you find a better cluster vs diagnoses match by cutting into a different number of clusters between 2 and 10? How do you judge the quality of your result in each case? Skipped.

Using Different Methods


Q12. Which method gives your favorite results for the same data.dist dataset? Explain your reasoning. My favorite is the ward.d2 method. I like the ward.d2

message because its easier to interpret.

Combining methods

Create a Ward clustering

```
d<-dist(wisc.pr$x [, 1:3])
wisc.pr.hclust <- hclust(d, method = "ward.D2")
plot(wisc.pr.hclust)
```



```
d
hclust (*, "ward.D2")
```


```
grps <- cutree(wisc.pr.hclust, k=2)
table(grps)
```

```
grps
 1 2
203 366
```


```
table(grps, diagnosis)
```

```
diagnosis
grps   B    M
1    24 179
2   333  33
```

```
ggplot(wisc.pr$x, aes(x = PC1, y = PC2, color = grps)) +
  geom_point(size = 2) +
  labs(x = "PC1", y = "PC2") +
  theme_minimal()
```



```
ggplot(wisc.pr$x, aes(x = PC1, y = PC2, color = diagnosis)) +
  geom_point(size = 2) +
  labs(x = "PC1", y = "PC2") +
  theme_minimal()
```



```
g <- as.factor(grps)
levels(g)
```


```
[1] "1" "2"
```

```
g <- relevel(g,2)
levels(g)
```

```
[1] "2" "1"
```

Plot your reordered factor:

```
ggplot(wisc.pr$x, aes(x = PC1, y = PC2, color = g)) +
  geom_point(size = 2) +
  labs(x = "PC1", y = "PC2") +
  theme_minimal()
```



I skipped the rgl and plotly packages.

```
wisc.pr.hclust <- hclust(dist(wisc.pr$x[, 1:7]), method = "ward.D2")
wisc.pr.hclust.clusters <- cutree(wisc.pr.hclust, k=2)
table(wisc.pr.hclust.clusters, diagnosis)
```

| | | diagnosis | |
|-------------------------|-----|-----------|---|
| wisc.pr.hclust.clusters | | B | M |
| 1 | 28 | 188 | |
| 2 | 329 | 24 | |

Q13. How well does the newly created model with four clusters separate out the two diagnoses? It does a really good job at separating the two diagnoses. I can easily tell that group 1 associates with mostly malignant whereas group 2 mostly associates with benign tumors.

```
table(wisc.hclust.clusters, wisc.df$diagnosis)
```

| wisc.hclust.clusters | | B | M |
|----------------------|----|-----|---|
| 1 | 12 | 165 | |
| 2 | 2 | 5 | |

| | | |
|---|-----|----|
| 3 | 343 | 40 |
| 4 | 0 | 2 |

Q14. How well do the hierarchical clustering models you created in previous sections (i.e. before PCA) do in terms of separating the diagnoses? Again, use the table() function to compare the output of each model (wisc.km\$cluster and wisc.hclust.clusters) with the vector containing the actual diagnoses. In the previous sections, the hierarchical models we're hard to interpret. The method that uses PCA is easier to interpret to make accurate predictions.

Sensitivity/Specificity

Q15. OPTIONAL: Which of your analysis procedures resulted in a clustering model with the best specificity? How about sensitivity? Skipped.

Prediction

```
#url <- "new_samples.csv"
url <- "https://tinyurl.com/new-samples-CSV"
new <- read.csv(url)
npc <- predict(wisc.pr, newdata=new)
npc
```

| | PC1 | PC2 | PC3 | PC4 | PC5 | PC6 | PC7 |
|------|--------------|-------------|--------------|--------------|-------------|--------------|------------|
| [1,] | 2.576616 | -3.135913 | 1.3990492 | -0.7631950 | 2.781648 | -0.8150185 | -0.3959098 |
| [2,] | -4.754928 | -3.009033 | -0.1660946 | -0.6052952 | -1.140698 | -1.2189945 | 0.8193031 |
| | PC8 | PC9 | PC10 | PC11 | PC12 | PC13 | PC14 |
| [1,] | -0.2307350 | 0.1029569 | -0.9272861 | 0.3411457 | 0.375921 | 0.1610764 | 1.187882 |
| [2,] | -0.3307423 | 0.5281896 | -0.4855301 | 0.7173233 | -1.185917 | 0.5893856 | 0.303029 |
| | PC15 | PC16 | PC17 | PC18 | PC19 | PC20 | |
| [1,] | 0.3216974 | -0.1743616 | -0.07875393 | -0.11207028 | -0.08802955 | -0.2495216 | |
| [2,] | 0.1299153 | 0.1448061 | -0.40509706 | 0.06565549 | 0.25591230 | -0.4289500 | |
| | PC21 | PC22 | PC23 | PC24 | PC25 | PC26 | |
| [1,] | 0.1228233 | 0.09358453 | 0.08347651 | 0.1223396 | 0.02124121 | 0.078884581 | |
| [2,] | -0.1224776 | 0.01732146 | 0.06316631 | -0.2338618 | -0.20755948 | -0.009833238 | |
| | PC27 | PC28 | PC29 | PC30 | | | |
| [1,] | 0.220199544 | -0.02946023 | -0.015620933 | 0.005269029 | | | |
| [2,] | -0.001134152 | 0.09638361 | 0.002795349 | -0.019015820 | | | |

```


library(ggplot2)

# Create a data frame for the PCA points
pca_df <- data.frame(PC1 = wisc.pr$x[, 1],
                      PC2 = wisc.pr$x[, 2],
                      group = as.factor(g))

# Create a data frame for npc points
npc_df <- data.frame(PC1 = npc[, 1],
                      PC2 = npc[, 2],
                      label = as.factor(c(1, 2)))

# Plot with ggplot2
ggplot(pca_df, aes(x = PC1, y = PC2, color = group)) +
  geom_point(alpha = 0.7) +
  geom_point(data = npc_df, aes(x = PC1, y = PC2),
             color = "blue", size = 6) +
  geom_text(data = npc_df, aes(label = label),
            color = "white", size = 4) +
  labs(x = "PC1",
       y = "PC2",
       title = "PCA Plot with Highlighted NPC Points") +
  theme_minimal()

```



Q16. Which of these new patients should we prioritize for follow up based on your results? Patient 2 needs to be prioritized for a check up since their tumor most likely will be malignant (since they fall into group 1). For group 1: TP=188 and FP=28